

Porting DOOM to Vulkan

SIGGRAPH 2016

Axel Gneiting

id Software



SIGGRAPH 2016



Agenda

- Demo & short idTech 6 overview
- Porting to Vulkan
 - Shaders, pipelines & states
 - Descriptor Sets
 - Multithreading
 - Image layouts & barriers
 - Memory & synchronization
 - Asynchronous compute
- Results & Future Work



DOOM



Video



SIGGRAPH 2016



idTech 6

- PC OpenGL & Vulkan, PS4, Xbox One
- DOOM and future id Software titles
- 60+ Hz on all Platforms
- Shader syntax similar to HLSL
 - Translated to PSSL/HLSL/GLSL at build time



CPU

- Parallel command buffer generation
 - Split up into several “contexts” per frame
 - Each contexts owns command buffer
 - For each context we run multiple jobs to fill CB
 - Last job in frame submits command buffers to GPU
- OpenGL runs sequential on one thread
 - Some scene preparation work is still in jobs



SIGGRAPH 2016



GPU

- Clustered forward shading with some deferred
- Same shader for most of the geometry
 - Same set of textures too (virtual texturing)
 - Very few state changes
- Extensive post process
 - DoF, Temporal AA, SSDO, motion blur, etc.
- Lots of asynchronous compute
 - DXT encode, particles & post processing



Porting to Vulkan

- Started 2015 with an early version
 - Wrote most of the Vulkan backend code
 - Got first triangle rendering
- Picked it up in late March 2016 again
- Was mostly running at game launch
 - RenderDoc helps, even better now!
- Small issues delaying release ☹
 - Driver issues
 - Swap chain surprisingly hard to get right



Porting to Vulkan

- Validation layers were unreliable back then
- Lots of false errors
- Had to write some validation code ourselves
- Validation layers much better now
- Still good to have own validation for debugging



Shaders

- Already had GLSL translator
 - But OpenGL was binding by name
 - Vulkan uses binding IDs at pipeline creation
- Using AMD extensions if available
 - Variant for all shaders
 - AMD_shader_ballot & AMD_gcn_shader



Shaders

- Normalized clip space is upside down
 - Shader generator adds `gl_Position.y = -gl_Position.y` at end of every vertex program
 - Can we please have an extension that fixes this?
 - Platform differences are a waste of time
- Z range is good: [0,1] ☺



Pipelines & States

- Abstraction layer still old style API like
- Need to emulate stateful API & track states
- Hash table for pipelines, render passes & frame buffer states
 - Way smaller perf overhead than thought
- Dynamic state for scissor/viewport/stencil and depth bias
- Only ~350 total graphics pipelines for entire game



Pipelines & States

- Pipeline creation expensive
 - Lookup misses unacceptable at runtime
 - Some pipelines take 100+ ms to compile
- Solution
 - Play game and serialize states to disk
 - On startup launch jobs to compile pipelines
 - Fairly robust, missed pipelines would just cause stalls for player



Descriptor Sets

- No deletion of Vulkan objects while playing
 - Geometry statically loaded
 - Textures virtualized
- Got away with a descriptor hash table
- One big descriptor set for each combination
- Complete table flush if a Vulkan handle gets deleted
 - Level load & unload, etc.
- About 3-4k descriptor sets usually



Descriptor Sets

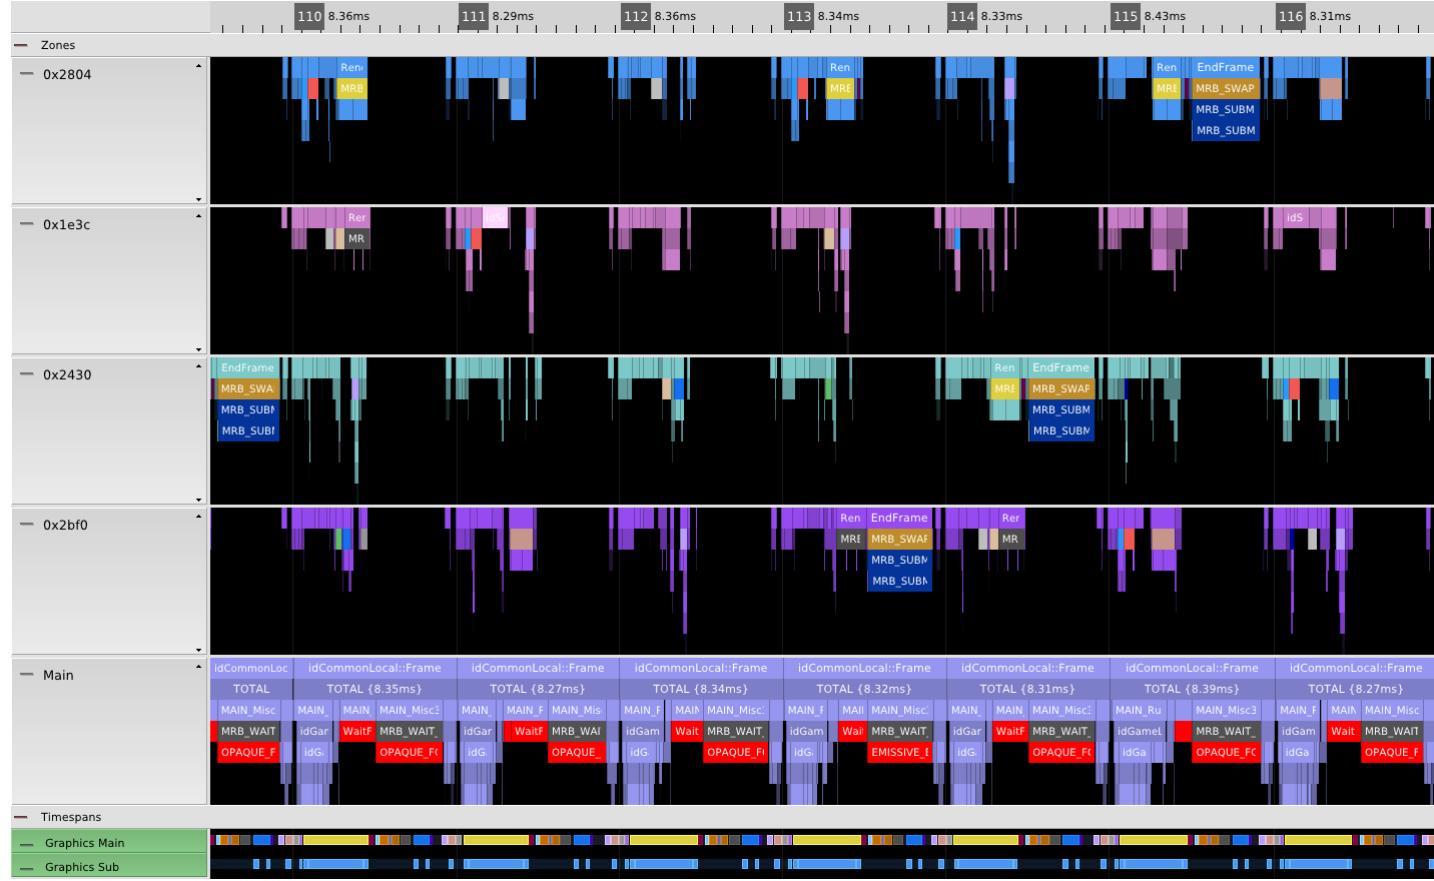
- Dynamic uniforms written to ring buffer
- Thread safe allocation from ring with atomics
 - 256 byte align allocations for simplicity
- Bound with UNIFORM_BUFFER_DYNAMIC
 - Offset set as vkCmdBindDescriptorSets parameter
- Also used UNIFORM_BUFFER_DYNAMIC for skinning data
 - Baked range problematic
 - Got away with 64kB range for everything
 - Alternative would have been way more descriptor sets



Multithreading

- Mostly straight forward port from consoles
- Image layouts problematic (more soon)
- Double buffered CBs per context
- Read/write locks for state hash tables
 - Never blocks if no state misses





SIGGRAPH 2016



Image layouts & barriers

- Image layouts were a big headache
 - 25+ barriers per frame
 - Hundreds of layout changes
- Combining as many barriers as possible
- Knowing last image state difficult
 - We only specify the new state in code
- But parallelism makes complete automatic tracking impossible



SIGGRAPH 2016

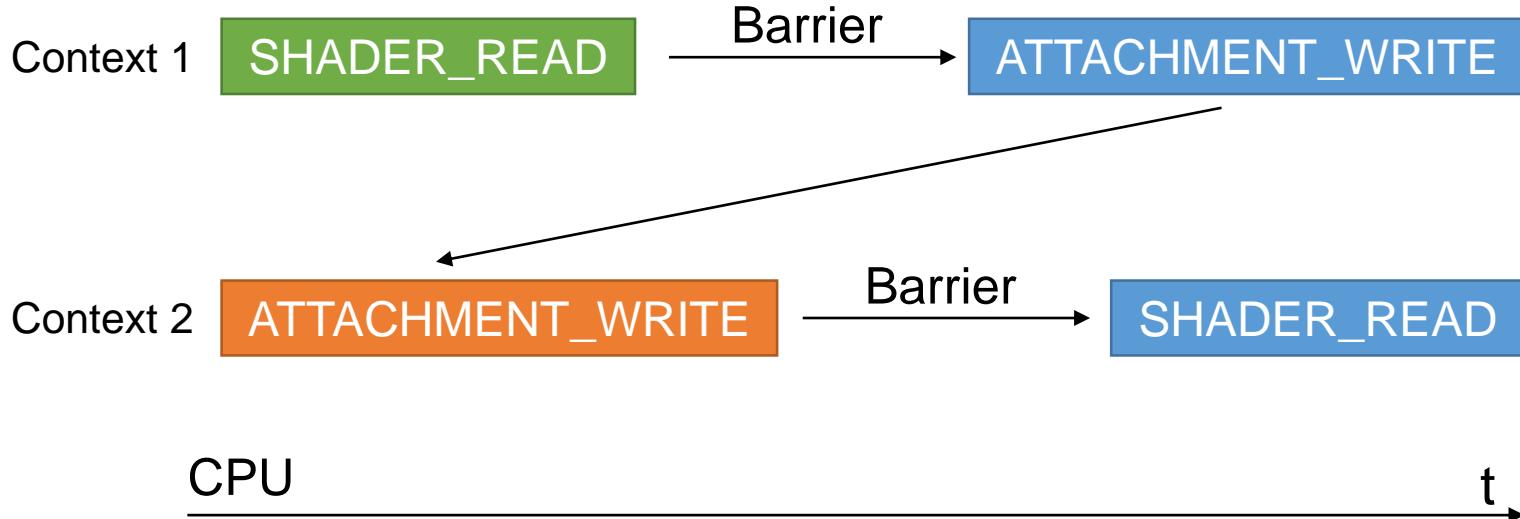


Image layouts & barriers

- Automatic tracking inside each context / CB
- Not many images used across CBs
- Start of frame: Set state for start of CB to fix up missing tracking
- End of frame:
 - Go over transitions & determine initial next frame state
 - Validate image transitions
- No vkCmdSetEvent/vkCmdWaitEvents right now



Image layouts & barriers



Memory

- Simple block allocator
 - Split into max 128 MB pieces
 - Try smaller allocation until allocation succeeds
 - Or falls back to system memory if allocations fail in VRAM
 - Resizable images allocated individually
- NVIDIA problematic under pressure (2GB)
 - Lots of fixes in driver by now
 - Use NV_dedicated_allocation if possible



Memory

- All uploads through common manager
- Double buffered host staging memory
- Each staging buffer associated with
 - Command buffer
 - Fence
- If buffer is full, write fence at end of CB and submit
- Wait on fence before reuse
- Flush host visible ranges before graphics submits



Synchronization

- Double buffering everywhere
 - Wait for command buffer fence on CPU
 - Minimizes latency
- GPUView is your friend!
 - Much more useful than with OpenGL/DX11
- Swap chains are tricky
 - Make sure acquire & present always matching
 - Acquire as late as possible (avoids stalls)



- Semaphore Wait
- Semaphore Signal
- Present
- Work (Submit)
- API Calls

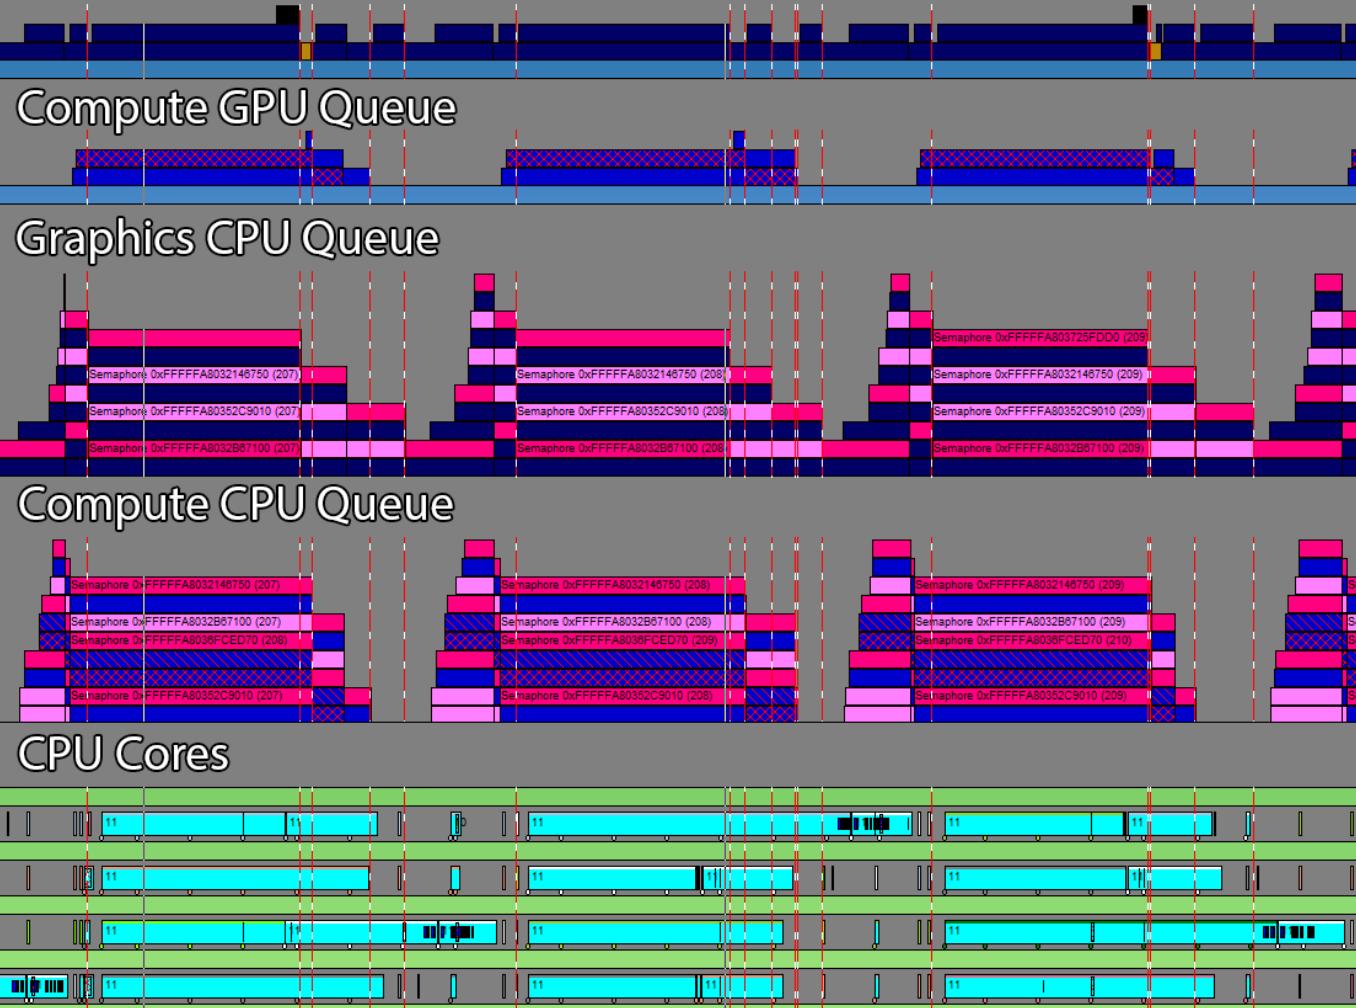
Graphics GPU Queue

Compute GPU Queue

Graphics CPU Queue

Compute CPU Queue

CPU Cores



Asynchronous Compute

- Useful for leveraging wasted GPU idle time
 - E.g. during shadow & depth pass
- GPU particles & post process
- Post process overlaps with beginning of next frame
 - Present from compute queue on AMD
 - NVIDIA still working on driver support
- Using SHARING_MODE_CONCURRENT for render targets
 - Careful, might be slower



Results

- Very pleased with performance gains
- 60%-70% in some scenes on AMD in GPU limit
 - Faster than OpenGL even without async/intrinsics
- NVIDIA GPU time about the same
- Render CPU limit is mostly gone
 - People reporting 60+ Hz in power saving mode
- Lots of potential



Future Work

- Prepare image barriers & layouts at beginning of frame
- Remove hashes and make high level code aware of states
- Know exactly what pipelines are used in game
- Better use of render passes (sub passes, layout transitions)



SIGGRAPH 2016



Future Work

- Split barriers (`vkCmdSetEvent/vkCmdWaitEvents`)
- Command buffer reuse (e.g. deferred passes & post process)
- More asynchronous compute
- Asynchronous transfers



Thanks

- Jean Geffroy, Tiago Sousa, Billy Khan & the whole team at id Software
- Baldur Karlsson for RenderDoc
- AMD and NVIDIA for help on Vulkan port
- Make sure to play the game!



We are Hiring

- Various openings across Zenimax Studios !
- Please visit <https://jobs.zenimax.com>

