




# Vulkan Subgroups

Nuno Subtil  
[nsubtil@nvidia.com](mailto:nsubtil@nvidia.com)

# Agenda

- What are subgroups and why they're useful
- Subgroup overview
- Vulkan 1.1 Subgroup Operations
- Partitioned Subgroup Operations
- NV Implementation Details
- Tips
- Mapping to HLSL


# Subgroups?



## Vulkan 1.0

- Threads execute in workgroups
- Each workgroup has some amount of (fast) shared memory
- Threads communicate via shared memory within workgroup

# Subgroups!



## Vulkan 1.0

- Threads execute in workgroups
- Each workgroup has some amount of (fast) shared memory
- Threads communicate via shared memory within workgroup

## Vulkan 1.1


- Adds **subgroups**: sets of threads within workgroup that can communicate directly
- Can be faster than shared memory

# Subgroups!

- All-to-all communication across sets of threads within a workgroup
  - Equivalent to HLSL SM6 wave ops
- Can be more efficient than shared memory
  - If the data you want to move is in registers, subgroups are typically faster
  - Implicit, finer-grained synchronization
- May have lower latency
  - Shared memory implies workgroup-wide synchronization
  - Subgroup operations only require synchronization within participating threads
- Exposed in all shader stages
  - Compute support is required
  - Other stages are allowed depending on the implementation

# Subgroup example: reduction


Simple parallel reduction example: sum of values across threads



# Subgroup example: reduction

Parallel reduction loop using shared memory


```
shared int s[WORKGROUP_SIZE];
...
int a = compute_local_value();
s[gl_WorkGroupID.x] = a;                                // memory write
memoryBarrierShared();                                    // synchronize workgroup
if (current_thread_should_reduce()) {
    int b = s[gl_WorkGroupID.x + iterDelta]; // memory read
    perform_reduction_step(a, b);
    iterDelta /= 2;
}
```



# Subgroup example: reduction

## Parallel reduction loop using shared memory


```
shared int s[WORKGROUP_SIZE];
...
int a = compute_local_value();
s[gl_WorkGroupID.x] = a;                                // memory write
memoryBarrierShared();                                    // synchronize workgroup
if (current_thread_should_reduce()) {
    int b = s[gl_WorkGroupID.x + iterDelta]; // memory read
    perform_reduction_step(a, b);
    iterDelta /= 2;
}
```



## Parallel reduction loop using subgroups

```
int a = compute_local_value();
int b = subgroupShuffleDown(a, iterDelta); // synchronize subgroup
if (current_thread_should_reduce()) {
    perform_reduction_step(a, b);
    iterDelta /= 2;
}
```


**ShuffleDown (x, 1)**



# Subgroup example: reduction

## Parallel reduction loop using shared memory

```
shared int s[WORKGROUP_SIZE];
...
int a = compute_local_value();
s[gl_WorkGroupID.x] = a;                                // memory write
memoryBarrierShared();                                    // synchronize workgroup
if (current_thread_should_reduce()) {
    int b = s[gl_WorkGroupID.x + iterDelta]; // memory read
    perform_reduction_step(a, b);
    iterDelta /= 2;
}
```



Workgroup-wide synchronization + two memory operations

## Parallel reduction loop using subgroups

```
int a = compute_local_value();
int b = subgroupShuffleDown(a, iterDelta); // synchronize subgroup
if (current_thread_should_reduce()) {
    perform_reduction_step(a, b);
    iterDelta /= 2;
}
```

Fewer threads synchronize; no memory operations

# When to use subgroups?

- Quite a few algorithms may benefit
  - Reductions: post-processing effects
    - Min/max/sum across range of data
    - Bloom, depth-of-field, motion blur are good candidates
  - List building: light culling
    - Reduce shared memory atomics, skip work across subgroup (e.g., entire subgroup decides no elements need to be added to the list)
  - Sorting
    - Bitonic sort can be accelerated via subgroups

# When *not* to use subgroups?

- Tradeoff between different kinds of latency
  - Shared memory:
    - Workgroup-wide synchronization latency
    - Read/write latency (e.g., if backed by cache)
  - Subgroups:
    - Subgroup-wide synchronization latency
    - Potentially increased ALU/issue latency (may expand to multiple instructions)
- Tradeoffs can be implementation dependent

# More details

- Set of shader invocations (threads)
  - *Efficiently* synchronize and share data with each other
  - Exposed “as if” running concurrently
    - Maps to warp (NV), wavefront (AMD)
    - Implementation *can* advertise smaller subgroup size than HW implements
- Invocations in a subgroup may be *active* or *inactive*
  - Active: execution is being performed
  - Inactive: not being executed
    - Non-uniform flow control
    - Local workgroup size not a multiple of subgroup size
  - Can change throughout shader execution as control flow diverges and re-converges

# Vulkan 1.1 API: Subgroup Properties

- A new structure to query subgroup support on a physical device
  - **subgroupSize** - number of invocations in each subgroup
    - *must be at least 1 (and <= 128)*
  - **supportedStages** - which shader stages support subgroup operations
    - **VK\_SHADER\_STAGE\_COMPUTE\_BIT is required**
  - **supportedOperations** - which subgroup operations are supported
    - **VK\_SUBGROUP\_FEATURE\_BASIC\_BIT is required**
  - **quadOperationsInAllStages** - do quads ops work in all stages or only fragment and compute

```
typedef struct VkPhysicalDeviceSubgroupProperties {  
    VkStructureType sType;  
    void* pNext;  
    uint32_t subgroupSize;  
    VkShaderStageFlags supportedStages;  
    VkSubgroupFeatureFlags supportedOperations;  
    VkBool32 quadOperationsInAllStages;  
} VkPhysicalDeviceSubgroupProperties;  
  
typedef enum VkSubgroupFeatureFlagBits {  
    VK_SUBGROUP_FEATURE_BASIC_BIT = 0x00000001,  
    VK_SUBGROUP_FEATURE_VOTE_BIT = 0x00000002,  
    VK_SUBGROUP_FEATURE_ARITHMETIC_BIT = 0x00000004,  
    VK_SUBGROUP_FEATURE_BALLOT_BIT = 0x00000008,  
    VK_SUBGROUP_FEATURE_SHUFFLE_BIT = 0x00000010,  
    VK_SUBGROUP_FEATURE_SHUFFLE_RELATIVE_BIT = 0x00000020,  
    VK_SUBGROUP_FEATURE_CLUSTERED_BIT = 0x00000040,  
    VK_SUBGROUP_FEATURE_QUAD_BIT = 0x00000080,  
    VK_SUBGROUP_FEATURE_PARTITIONED_BIT_NV = 0x00000100,  
} VkSubgroupFeatureFlagBits;
```

# Shader Built-in variables

- All supported stages
  - `gl_SubgroupSize` - size of the subgroup - matches the API property
  - `gl_SubgroupInvocationID` - ID of the invocation within the subgroup, [0..`gl_SubgroupSize`)
  - `gl_Subgroup{Eq,Ge,Gt,Le,Lt}Mask`
    - bitmask of all invocations as `compared` to the `gl_SubgroupInvocationID` of current invocation
    - Useful for working with subgroupBallot results (more on this later)
- Compute only
  - `gl_NumSubgroups` - number of subgroups in local workgroup
  - `gl_SubgroupID` - ID of subgroup within local workgroup, [0..`gl_NumSubgroups`)

	T	T	T	T		T	T	T	T
<code>gl_SubgroupSize:</code>	4	4	4	4		4	4	4	4
<code>gl_SubgroupInvocationID:</code>	0	1	2	3		0	1	2	3
<code>gl_SubgroupID:</code>	0	0	0	0		1	1	1	1
<code>gl_SubgroupLtMask:</code>	0	1	3	7		0	1	3	7


# Subgroup Basic Operations

- Subgroup-wide barriers

- void `subgroupBarrier()` - Full memory and execution barrier
  - All active `invocations` `sync` and `memory stores` to *coherent* memory locations are `completed`
- void `subgroupMemoryBarrier()`
  - Enforces ordering of `all memory` transactions by an invocation, as seen by other invocations in the subgroup
- void `subgroupMemoryBarrier{Buffer,Image,Shared}()`
  - Enforces ordering on `buffer`, `image`, or `shared` (compute only) `memory` operations, respectively

# Subgroup Vote Operations

- Select one thread in a subgroup
  - bool `subgroupElect()`
    - Pick one active invocation, always the one with lowest `gl_SubgroupInvocationID`
    - Used for executing work on only one invocation



# Subgroup Vote Operations


- Determine if a Boolean condition is met **across the entire subgroup**
  - bool `subgroupAll`(bool value)
    - true if **all** invocations have <value> == true
  - bool `subgroupAny`(bool value)
    - true if **any** invocation has <value> == true
  - bool `subgroupAllEqual`(T value)
    - true if **all** invocations have the **same value** of <value>
- Useful for code that has branching
  - Can do more optimal calculations if certain conditions are met

```
void main() {
    bool condition = foo[gl_GlobalInvocationID.x] < bar[gl_GlobalInvocationID.x];

    if (subgroupAll(condition)) {
        // all invocations in the subgroup are performing x
    } else if (!subgroupAny(condition)) {
        // all invocations in the subgroup are performing y
    } else {
        // Invocations that get here are doing a mix of x & y so have a fallback
    }
}
```

# Subgroup Arithmetic Operations


- Operations across **all** active **invocations** in a subgroup
  - $T \text{ subgroup} <\text{op}> (T \text{ value})$ 
    - $<\text{op}>$  = **Add**, **Mul**, **Min**, **Max**, **And**, **Or**, **Xor**
  - Reduction operations
    - Returns the result of the **same calculation** to each invocation




# Subgroup Arithmetic Operations

- Operation on **invocations** with `gl_SubgroupInvocationID` **less than self**
  - `T subgroupInclusive<op>(T value)`
    - Includes **own** value in operation
  - `T subgroupExclusive<op>(T values)`
    - Excludes **own** value from operation
  - Inclusive or exclusive scan
    - Returns the result of **different calculation** to each invocation

InclusiveAdd




ExclusiveAdd




# Subgroup Ballot Operations

- Allow invocations to do limited **data sharing** across a subgroup
  - **Broadcast** - value from one invocation to all invocations
    - `T subgroupBroadcast(T value, uint id)`
      - Broadcasts `<value>` from the invocation with `gl_SubgroupInvocationID == id`
      - `<id>` must be compile time constant
    - `T subgroupBroadcastFirst(T value)`
      - Broadcasts `<value>` from the invocation with `lowest active gl_SubgroupInvocationID`

**Broadcast(x,2)**



**BroadcastFirst**



# Subgroup Ballot Operations

- Allow invocations to do limited **data sharing** across a subgroup
  - More powerful form of voting
    - `uvec4 subgroupBallot(bool value)`
      - Returns bitfield **ballot** with result of evaluating <value> in each invocation
      - Bit <i> == 1 means expression evaluated to true for `gl_SubgroupInvocationID == i`
      - Bit <i> == 0 means expression evaluated to false, or invocation inactive
      - `uvec4` used in ballots is treated as a bitfield with `gl_SubgroupSize` significant bits
      - First invocation is in bit 0 of first vector component (.x), 32<sup>nd</sup> invocation in bit 0 of .y, etc.
      - Bits beyond `gl_SubgroupSize` are ignored
      - `subgroupBallot(true)` gives a bitfield of all the active invocations

4	5	1	7
---	---	---	---

**Ballot**(val > 4) → 0b1010

4	5	1	7
---	---	---	---

**Ballot**(true) → 0b1111

5	7
---	---

**Ballot**(true) → 0b1010

# Subgroup Ballot Operations

- Ballot helper functions - to simplify working with uvec4 bitfield
  - bool `subgroupInverseBallot`(uvec4 value)
    - Returns true if current invocation bit in <value> is set
  - bool `subgroupBallotBitExtract`(uvec4 value, uint index)
    - Returns true if bit in <value> that corresponds to <index> is 1
  - uint `subgroupBallotBitCount`(uvec4 value)
    - Returns the count of bits set to 1 in <value>
  - uint `subgroupBallot{Exclusive,Inclusive}BitCount`(uvec4 value)
    - Returns the exclusive/inclusive count of bits set to 1 in <value>
    - For bits with invocation ID < or <= the current invocation ID
  - uint `subgroupBallotFind{LSB,MSB}`(uvec4 value)
    - Returns the lowest/highest bit set in <value>

**InverseBallot(0b1010)** → 

F	T	F	T
---	---	---	---

**BallotInclusiveBitCount(0b1101)**

→ 

1	1	2	3
---	---	---	---

$\text{== } \text{BallotBitCount}(\text{val} \& \text{gl\_SubgroupLeMask})$

**BallotBitCount(0b1101)** → 

3	3	3	3
---	---	---	---

**BallotExclusiveBitCount(0b1101)**

→ 


0	1	1	2
---	---	---	---

$\text{== } \text{BallotBitCount}(\text{val} \& \text{gl\_SubgroupLtMask})$

# Subgroup Shuffle Operations

- More extensive data sharing across the subgroup
  - Invocations can **read values from other invocations** in the subgroup
- **Shuffle**
  - T **subgroupShuffle(T value, uint id)**
    - Returns <value> from the invocation with gl\_SubgroupInvocationID == **id**
    - Like subgroupBroadcast, but <id> can be determined **dynamically**

**Shuffle(x, index)**  
index = 2, 1, 1, 0




# Subgroup Shuffle Operations


- **ShuffleXor**

- T `subgroupShuffleXor(T value, uint mask)`
  - Returns `<value>` from the invocation with `gl_SubgroupInvocationID == (mask ^ current)`
  - Every invocation trades value with **exactly one other** invocation
  - Specialization of general shuffle
  - `<mask>` must be constant integral expression
  - Special conditions for using in a loop (basically needs to be unrollable)

**ShuffleXor(x, 1)**




**ShuffleXor(x, 2)**




# Subgroup Shuffle Relative Operations

- Enable invocations to perform **shifted data sharing** between subgroup invocations
  - T **subgroupShuffleUp**(T value, uint delta)
    - Returns <value> from the invocation with gl\_SubgroupInvocationID == (**current - delta**)
  - T **subgroupShuffleDown**(T value, uint delta)
    - Returns <values> from the invocation with gl\_SubgroupInvocationID == (**current + delta**)
- Useful to construct your own scan operations
  - Strided scan (e.g. even or odd invocations, etc.)


ShuffleUp(x, 1)




ShuffleUp(x, 2)



ShuffleDown (x, 1)




ShuffleDown (x, 2)




# Subgroup Clustered Operations

- Perform arithmetic operations across a **fixed partition** of a subgroup
  - T `subgroupClustered<op>(T value, uint clusterSize)`
    - <op> = `Add`, `Mul`, `Min`, `Max`, `And`, `Or`, `Xor`
    - clusterSize - size of partition
      - compile-time constant
      - power of 2,  $\geq 1$
    - Only active invocations in the partition participate
  - Sharing data only with a selection of your closest neighbors
    - An algorithm that relies on a fixed size grid `< gl_SubgroupSize`
    - Eg: Convolution neural network - max pooling
      - Take large data set and compress to a smaller one
      - Divide data into  $N \times N$  grid -  $N = \text{clusterSize}$
      - Output maximum for each cluster

**ClusteredAdd(x, 2)**



**ClusteredMax(x, 4)**



# Subgroup Quad Operations


- Subgroup **quad** is a cluster of size 4
  - Neighboring pixels in a 2x2 grid in fragment shaders (ie derivative group)
  - Not restricted to fragment shaders,
  - Just a cluster of 4 in other stages (no defined layout)
    - Remember to check for support (`quadOperationsInAllStages` property)

0	1
2	3

0	1	2	3
---	---	---	---

## • Broadcast



- T `subgroupQuadBroadcast(T value, uint id)`
  - Returns <value> from the invocation where `gl_SubgroupInvocationID % 4 = <id>`



# Subgroup Quad Operations


- Swap

- T `subgroupQuadSwapHorizontal(T value)`
  - Swap values **horizontally** in the quad
- T `subgroupQuadSwapVertical(T value)`
  - Swap values **vertically** in the quad
- T `subgroupQuadSwapDiagonal(T value)`
  - Swap values **diagonally** in the quad
- Can easily construct a lower resolution image (2x2 filter)
  - See subgroup tutorial for details



# Subgroup Partitioned Operations (NV)


- Perform arithmetic operations across a **flexible set** of invocations
  - Generalization of clustering which does not need fixed-size clusters or offsets
  - [VK\\_NV\\_shader\\_subgroup\\_partitioned](#) / [GL\\_NV\\_shader\\_subgroup\\_partitioned](#)
- Generate a partition
  - `uvec4 subgroupPartitionNV(T value)`
    - Returns a **ballot** which is a **partition** of all invocations in the subgroup based on <value>
    - All invocations represented by the same ballot have the same <value>
    - All invocations in different ballots have different <value>



# Subgroup Partitioned Operations (NV)

- Operation on a partition
  - T `subgroupPartitionedInclusive<op>NV`(T value, uvec4 ballot)
  - T `subgroupPartitionedExclusive<op>NV`(T value, uvec4 ballot)
  - T `subgroupPartitioned<op>NV`(T value, uvec4 ballot)
    - <op> is `Add`, `Mul`, `Min`, `Max`, `And`, `Or`, `Xor`
    - Inclusive scan, exclusive scan, reduction operate similar to clustered/arithmetic operations
    - <ballot> describes the partition - typically the result from `subgroupPartitionNV`
    - No restrictions on how the invocations are partitioned, except that the ballot values passed in must represent a “valid” partition

**PartitionedAddNV(values, ballot)**  
ballot = 0x5, 0x12, 0x28, 0xC0



# Subgroup Partitioned Operations

- Why partitions?
  - Shaders can't really predict that consecutive invocations will have related values
  - More useful to "discover" (subgroupPartitionNV) those invocations that are related, and then do subgroup operations on related invocations
  - E.g. Deferred shading, detect pixels with the same material or light
- Any implementation that supports `VK_SUBGROUP_FEATURE_ARITHMETIC_BIT` can trivially support partitioned ops
  - Loop over unique partition subsets, compute each in flow control
  - Cost = NumSubsets \* costof(SUBGROUP\_FEATURE\_ARITHMETIC)
- Some implementations can compute all subsets in parallel
  - Cost = costof(SUBGROUP\_FEATURE\_ARITHMETIC)
  - More useful generalization of clustering, and at the same cost
- Most implementations can probably do better than the trivial looping

# NVIDIA Implementation Details

- Native hw instructions are essentially what is exposed in
  - [GL\\_NV\\_shader\\_thread\\_shuffle](#) and [GL\\_NV\\_shader\\_thread\\_group](#)
- **shuffle/shuffleUp/shuffleDown/shuffleXor** are **fast** instructions
  - Essentially our primitives
  - Most other instructions are built up from these using relatively simple transforms
  - Don't be afraid to use more general operations!
    - Can still be faster than composing from building blocks
- All the subgroup operations are similar cost
  - E.g. a REDUCE operation (subgroup<op>) is basically:

```
x = op(x, shuffleXor(x, 1));  
x = op(x, shuffleXor(x, 2));  
x = op(x, shuffleXor(x, 4));  
x = op(x, shuffleXor(x, 8));  
x = op(x, shuffleXor(x, 16));
```

# Tips

- Make local workgroup be at least the size of the subgroup (compute),
  - Ideally integer multiples
  - Common subgroup sizes: 32 (NVIDIA, Intel), 64 (AMD)
- Subgroup size of 1 isn't very useful, but makes a single code path possible
- Subgroup operations provide implicit subgroup execution barriers
- Operations only act on **active** invocations
- Be aware of inactive lanes or out of range invocation IDs
  - Reading gives undefined values in most cases!
- Helper invocations participate in subgroup operations

# HLSL SM 6.0 Wave Ops Comparison

## D3D Wave Ops

- Wave lane count: 4 - 128
- Required in pixel and compute shaders
  - Not supported in any other stages
- All or nothing functionality
- Types: half, float, double, int, uint, short, ushort, uint64 (as supported)

## Vulkan Subgroups

- Subgroup size: 1 - 128
- Required in compute shaders
  - Optional in Frag, Vert, Tess, Geom stages
- Minimum functionality guaranteed, additional bundles of functionality
- Types: bool, float, double, int, uint
  - More types to be added in the future
- More complete set of intrinsics
  - Inclusive scan, clustered ops, etc.
  - Barriers
  - More helper routines

# Availability

- GLSL functionality
  - Glslang - <https://github.com/khronosgroup/glslang/>
- HLSL functionality
  - Glslang - <https://github.com/KhronosGroup/glslang>
  - DXC - <https://github.com/Microsoft/DirectXShaderCompiler/>
- SPIR-V 1.3
- Vulkan support
  - <https://vulkan.gpuinfo.org/> (under Device Properties)
  - NVIDIA Vulkan 1.1 drivers - <http://www.nvidia.com/Download/index.aspx>
  - AMD Vulkan 1.1 drivers
  - Intel Vulkan 1.1 drivers

# References

- Vulkan Subgroup Tutorial
  - <https://www.khronos.org/blog/vulkan-subgroup-tutorial>
- GL\_KHR\_shader\_subgroup GLSL extension
  - [https://github.com/KhronosGroup/GLSL/blob/master/extensions/khr/GL\\_KHR\\_shader\\_subgroup.txt](https://github.com/KhronosGroup/GLSL/blob/master/extensions/khr/GL_KHR_shader_subgroup.txt)
- GL\_NV\_shader\_subgroup\_partitioned GLSL extension
  - [https://github.com/KhronosGroup/GLSL/blob/master/extensions/nv/GL\\_NV\\_shader\\_subgroup\\_partitioned.txt](https://github.com/KhronosGroup/GLSL/blob/master/extensions/nv/GL_NV_shader_subgroup_partitioned.txt)
- HLSL Shader Model 6.0 (MSDN)
  - [https://msdn.microsoft.com/en-us/library/windows/desktop/mt733232\(v=vs.85\).aspx](https://msdn.microsoft.com/en-us/library/windows/desktop/mt733232(v=vs.85).aspx)
- DirectXShaderCompiler Wave Intrinsics
  - <https://github.com/Microsoft/DirectXShaderCompiler/wiki/Wave-Intrinsics>
- Reading Between the Threads: Shader Intrinsics
  - <https://developer.nvidia.com/reading-between-threads-shader-intrinsics>
- Faster Parallel Reductions on Kepler
  - <https://devblogs.nvidia.com/faster-parallel-reductions-kepler/>

# Thank You

- Daniel Koch (NVIDIA)
- Neil Henning (AMD) @sheredom
- Lei Zhang (Google)
- Jeff Bolz (NVIDIA)
- Piers Daniell (NVIDIA)

# HLSL / GLSL / SPIR-V Mappings

HLSL Intrinsic (Query)	GLSL Intrinsic	SPIR-V Op
WaveGetLaneCount() [4-128]	gl_SubgroupSize[1-128]	SubgroupSize decorated OpVariable
WaveGetLaneIndex	gl_SubgroupInvocationID	SubgroupId decorated OpVariable
WavesFirstLane()	subgroupElect()	OpGroupNonUniformElect

HLSL Intrinsic (Vote)	GLSL Intrinsic	SPIR-V Op
WaveActiveAnyTrue()	subgroupAny()	OpGroupNonUniformAny
WaveActiveAllTrue()	subgroupAll()	OpGroupNonUniformAll
WaveActiveBallot()	subgroupBallot()	OpGroupNonUniformBallot

HLSL Intrinsic (Broadcast)	GLSL Intrinsic	SPIR-V Op
WaveReadLaneAt()	subgroupBroadcast(const) / subgroupShuffle(dynamic)	OpGroupNonUniformBroadcast / OpGroupNonUniformShuffle
WaveReadLaneFirst()	subgroupBroadcastFirst()	OpGroupNonUniformBroadcastFirst

# HLSL / GLSL / SPIR-V Mappings

HLSL Intrinsic (Reduction)	GLSL Intrinsic	SPIR-V Op
WaveActiveAllEqual()	subgroupAllEqual()	OpGroupNonUniformAllEqual
WaveActiveBitAnd()	subgroupAnd()	OpGroupNonUniformBitwiseAnd / OpGroupNonUniformLogicalAnd
WaveActiveBitOr()	subgroupOr()	OpGroupNonUniformBitwiseOr / OpGroupNonUniformLogicalOr
WaveActiveBitXor()	subgroupXor()	OpGroupNonUniformBitwiseXor / OpGroupNonUniformLogicalXor
WaveActiveCountBits()	subgroupBallotBitcount()	OpGroupNonUniformBallotBitCount
WaveActiveMax()	subgroupMax()	OpGroupNonUniform*Max
WaveActiveMin()	subgroupMin()	OpGroupNonUniform*Min
WaveActiveProduct()	subgroupMul()	OpGroupNonUniform*Mul
WaveActiveSum()	subgroupAdd()	OpGroupNonUniform*Add

# HLSL / GLSL / SPIR-V Mappings

HLSL Intrinsic (Scan and Prefix)	GLSL Intrinsic	SPIR-V Op
WavePrefixCountBits()	subgroupBallotExclusiveBitCount()	OpGroupNonUniformBallotBitCount
WavePrefixSum()	subgroupExclusiveAdd()	OpGroupNonUniform*Add
WavePrefixProduct()	subgroupExclusiveMul()	OpGroupNonUniform*Mul

HLSL Intrinsic (Quad Shuffle)	GLSL Intrinsic	SPIR-V Op
QuadReadLaneAt()	subgroupQuadBroadcast()	OpGroupNonUniformQuadBroadcast
QuadReadAcrossDiagonal()	subgroupQuadSwapDiagonal()	OpGroupNonUniformQuadSwap
QuadReadAcrossX()	subgroupQuadSwapHorizontal()	OpGroupNonUniformQuadSwap
QuadReadAcrossY()	subgroupQuadSwapVertical()	OpGroupNonUniformQuadSwap